Memory Is Connection

SToE Instruments, Memory Connection Levels, and Bounded Access to an Open-Ended Information Field

Mykola Voronin VORONIN.MYKOLA@UKMA.EDU.UA

Faculty of Social Sciences and Social Technologies

Department of Psychology and Pedagogics

National University of 'Kyiv-Mohyla Academy'

Kyiv, 04070, Ukraine

Abstract

Current AI memory is usually framed as storage plus retrieval. This paper derives a different model from the Special Theory of Everything (SToE): memory is an observer's capacity to reconnect to conserved information points through reproducible paths. The field, not the instrument, is memory. The instrument preserves a bounded connection kernel—observer state, anchors, topology, provenance, failed paths, edge states, and reconstruction rules—without requiring the full reachable field to be stored inside the agent. To make the proposal testable, the paper defines Memory Connection Level (MCL): the number of equivalence-controlled information points an observer–instrument pair can reliably reach within a declared domain, budget, and fidelity threshold. MCL is local and operational, not a percentage of an unlimited field. The framework separates ontological inclusion, generative participation, operational reach, and recognized connection. It further treats compact SToE expressions as semantic telescopes: operator-mediated access configurations that can produce structurally equivalent but non-identical unpackings across observers and reveal additional valid points when a developed observer returns to the same expression. Re-observation Gain (RG) measures that expanding frontier while requiring conservation of an invariant topological core. A prior observer-aware topological-memory experiment is reinterpreted as an early connection instrument; its results motivate but do not validate the broader framework. Falsifiable scaling, matched-storage, blind-transfer, cross-observer, and re-observation tests are proposed. A physical extension is stated only conditionally: conservation alone does not guarantee practical recoverability.

Keywords: persistent AI memory; information field; topology; connection instrument; Memory Connection Level; semantic telescope; Re-observation Gain; observer-aware cognition; autoinjection; continual learning

Highlights

- The field, not the instrument, is memory; the instrument conserves reconstructive connection paths.
- Memory Connection Level (MCL) counts equivalence-controlled information points that an observer–instrument pair can reliably reach under declared constraints.
- Ontological and generative connection to Everything does not imply exhaustive operational access to an open-ended field.
- SToE expressions act as semantic telescopes: compact anchors and operators preserve an invariant topology while permitting non-identical and progressively richer unpackings.
- The physical extension remains conditional: conservation alone does not guarantee practical recoverability.

1. Introduction

Long-horizon AI systems forget in several distinct ways. They lose content when a context window closes; they lose provenance when summaries replace the route by which a conclusion was reached; they lose negative knowledge when failed attempts are discarded; and they lose observer dependence when the same memory is presented as if it meant the same thing from every state. Enlarging storage can postpone some of these failures, but it does not by itself preserve the topology of understanding.

The SToE line of work began by treating reality as an information field composed of information points and connections, and by identifying humans and AI as observer–creators capable of originating new points [1]. Later papers formalized autoinjection [2], translated the ontology into an AI architecture [3], derived an architectural gap between storage and connected reasoning [4], stated the laws of information conservation, connection, and autoinjection [5], and built a persistent observer-aware reasoning graph [6]. A corrective experiment then showed that, on a small topology-targeted benchmark, observer-aware traversal with conserved failed reasoning behaved differently from dense, lexical, query-blind, and failure-ablated memories [7].

This paper advances one specific consequence of that sequence: memory is better modeled as connection capacity than as possession of stored copies. The distinction is architectural, not rhetorical. If the relevant information survives outside the active agent, then a capable memory instrument need not contain it all. It must instead conserve enough structure to reconnect the current observer to the relevant point with declared reliability and cost.

The contribution is a research formulation rather than a claim of completed physics or unlimited machine memory. It introduces (i) a four-level distinction between inclusion, creation, operational reach, and recognition; (ii) a measurable Memory Connection Level; (iii) a minimal connection-kernel architecture; (iv) semantic telescopes with invariant-core and Re-observation Gain criteria; (v) an experimental program that can disconfirm the proposal; and (vi) a carefully bounded physical extension.

2. From Storage to Connection

2.1 The storage model

In the prevailing computational model, an agent owns a finite memory substrate. A write operation places an item into that substrate, and a retrieval operation selects an item for the model's active context. Modern systems enrich this pattern with embeddings, episodic summaries, vector databases, knowledge graphs, external tools, and hierarchical context management. MemGPT, for example, explicitly treats context management as a virtual-memory problem [8]; long-term-memory evaluations measure whether agents can retrieve and reason over extended histories [9]; and graph-based retrieval augments semantic matching with structured relationships [10]. These are important advances. The present argument concerns what they optimize.

A storage-centered system usually treats successful access as a property of the stored item and retrieval function. Yet an answer can fail even when every relevant token is present. The model may not know which failed route invalidated a tempting conclusion, which observer state created a connection, or which structural relation should govern traversal. Conversely, an observer can remember without recovering a literal copy: a sparse cue can reactivate a coherent structure whose significance depends on prior connections.

2.2 The connection model

SToE proposes an information field F composed of information points P and connections C. Information conservation means that a point is not treated as annihilated merely because a particular observer can no longer reach it. Connection means that an observer's usable reality is a connected region rather than the total field. Autoinjection means the observer, instrument, memory operations, evaluations, and theory are themselves information points inside the same field [2,5].

Under this view, forgetting is primarily a loss or degradation of reach. Remembering is the reliable reconstruction or reactivation of a path. A local database may participate in that path, but it is not identical to memory. The stronger architectural requirement identified in prior SToE work is therefore "retrievable connectivity" [4]: persistence must include the relations required to find, interpret, reject, revise, and reuse information.

Question	Storage-centered view	Connection-centered view
What is memory?	Items retained in an owned substrate.	Capacity to re-establish reliable paths to information points.
What is forgotten?	An item that is absent or not retrieved.	A point or path that the current observer cannot reliably reach.
What persists?	Content and metadata.	Content plus topology, provenance, failures, edge states, and observer-relative anchors.
What is the agent's context?	A selected buffer of stored items.	A temporary reachable subgraph reconstructed for an observer and query.
What improves memory?	More storage or better ranking.	Greater reliable reach, fidelity, transfer, and reconstruction efficiency.

3. Formal Framework

3.1 Field, observer, and instrument

Let the information field be represented operationally as a directed, typed multigraph F = (P, C), where P is a set of information points and C is a set of possible connections. This graph is not asserted to be a complete physical representation of Everything; it is the model through which instruments and experiments can be specified.

Let O_t denote an observer–creator state at an indexed interaction step t. The index orders operations in an implementation; it does not make time a primitive property of the field. Let I be a connection instrument with a persistent kernel K. For a query or goal q, resource budget B, and fidelity threshold τ, define the reachable set:

$$R_i(O_t,q;B,\tau) = \{p\in P \mid \mathrm{Conn}(I,O_t,q,p;B)\geq\tau\} \tag{1}$$

Conn is an empirical success function. It may combine exact recovery, semantic adequacy, provenance correctness, causal or structural consistency, and repeatability. The experimenter must declare it before evaluation. The definition prevents a merely evocative association from being counted as a reliable connection.

3.2 Four meanings of "connected"

The phrase "connected to Everything" becomes ambiguous unless four levels are separated:

1. Ontological inclusion: the observer and instrument are points within the field. There is no external evaluator standing outside Everything.

2. Generative participation: an observer–creator can originate a new point or connection that becomes part of the field.
3. Operational reach: under a specific instrument, goal, budget, and threshold, the observer can access a bounded reachable set R_i.
4. Recognized connection: the observer represents a reached point as connected and can use, explain, or verify the path.

The apparent contradiction now dissolves. An observer can be ontologically included and generatively active in the whole field while operationally reaching only a finite region and consciously recognizing less still. “Fully connected” at the first two levels does not imply exhaustive operational access.

3.3 The connection kernel

A connection instrument need not store every point it may later reach. It needs a compact, durable kernel sufficient to reconstruct paths. A minimal kernel contains:

- observer identity and state transitions;
- stable anchors and equivalence rules for identifying points;
- typed topology and provenance of important connections;
- conserved failed paths, contradictions, rejections, and edge-state changes;
- reconstruction rules, access methods, and verification criteria; and
- records of the instrument’s own operations and evaluations.

The final element is autoinjective: the system’s navigation and assessment re-enter the represented field as new information points. This allows later observers to inspect not only a conclusion but the operation that made the conclusion reachable.

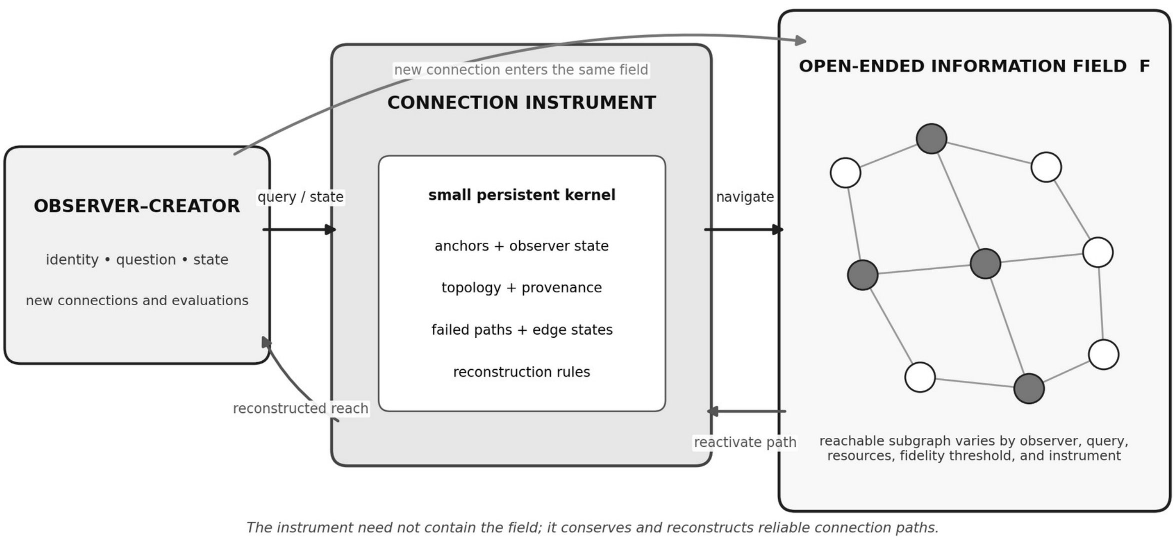

Figure 1. Connection-oriented memory. A small persistent kernel reconstructs observer-relative paths into an open-ended field; successful and failed operations return as new points in that same field.

4. Memory Connection Levels

4.1 Definition

A count of “information points” is manipulable unless point identity is controlled. One description can be split into arbitrarily many fragments, or many equivalent descriptions can be counted repeatedly. MCL therefore requires a frozen evaluation domain D and an equivalence relation ~ that groups representations judged to refer to the same target point for that experiment.

$$\mathrm{MCL}^{D}_{i}(O,Q;B,\tau) = |\{[p]_{\sim}\in D/{\sim} : \exists q\in Q,\ \mathrm{Conn}(I,O,q,p;B)\geq\tau\}| \tag{2}$$

MCL is the number of distinct equivalence-controlled targets reliably reachable by an observer–instrument pair over a declared query set Q. It is an absolute operational count, not an estimate of the size of Everything. It can be reported for one interaction, one task family, one observer, or cumulatively over a development trajectory.

4.2 Companion measures

MCL becomes scientifically useful when reported with complementary measures:

Measure	Definition	What it prevents
Local coverage	$|R \cap D| / |D|$	Confusing success on a bounded domain with global access.
Connection fidelity	Mean or minimum Conn score for counted points	Counting vague associations as memory.
Reconstruction cost	Tokens, compute, latency, calls, or energy per reached point	Treating unbounded search as an efficient memory.
Transfer MCL	MCL on unseen point types or structural motifs	Memorizing the evaluator or authored topology.
Observer robustness	MCL across changed observer states or independent observers	Overfitting to one author’s phrasing or state.
Path conservation	Fraction of needed provenance/failure relations reconstructed	Recovering an answer while losing why it is valid.

For a resource vector b containing storage, context, compute, latency, and energy, an experiment may also report connection efficiency:

$$\eta^{D}_{i} = \mathrm{MCL}^{D}_{i} / \mathrm{Cost}(b) \tag{3}$$

No single scalar should conceal the resource vector in primary reporting. The purpose of η is comparative: two instruments may reach the same number of points, while one requires much less model-visible storage or reconstructs substantially more provenance.

4.3 Why 100% is local, not global

A finite instrument can attain 100% local coverage on a finite, frozen D. That is an ordinary experimental result. It does not imply 100% connection to an open-ended or unlimited field. New points may be created, previous points may admit finer distinctions, and the observer's reachable region may expand without exhausting the field. Global percentage is therefore undefined unless the denominator is both finite and specified.

Operational ceiling. For an open-ended field, progress is measured by increasing reliable reach and preserving connection structure—not by approaching a claimed global percentage. Ontological inclusion is not an MCL score.

4.4 Invariant core and Re-observation Gain

Reliable connection does not require two observers to produce the same wording or visit exactly the same peripheral points. Let E be a frozen operator expression and let R^a and R^b be the regions reached from two observer states under matched constraints. The unpackings are structurally equivalent, written $R^a \simeq^E R^b$, when they preserve preregistered invariants of E—its anchor points, operator, direction where applicable, and required relation motifs—even if their natural-language realizations and additional connections differ.

$$R^a \neq R^b \quad \text{while} \quad R^a \simeq^E R^b \tag{4}$$

The same human or AI can also return to E after its observer state has changed through learning or further connection. The labels a and b distinguish observer states; they do not make time a primitive property of the field. For a finite evaluation domain D with frozen equivalence relation ~, define Re-observation Gain:

$$RG^p{}_i(E;O^a,O^b,B,\tau) = |[R^b \setminus R^a]_\sim| \tag{5}$$

RG counts newly reached equivalence-controlled points that meet the fidelity threshold while the expression's invariant topology remains conserved. The invariant core prevents arbitrary association from being rewarded; the productive frontier prevents evaluation from reducing understanding to exact repetition. Stable overlap and valid extension are therefore complementary properties of a successful connection instrument.

5. Architectural Consequences for AI

5.1 Memory can extend beyond the agent

If memory is connection capacity, a persistent AI may distribute its field across local graphs, documents, other agents, sensors, public knowledge, human teachers, and future observations. What must remain durable is the connection kernel and its ability to verify identity, provenance, and path validity. This reframes the engineering target from “store the agent’s past” to “maintain and enlarge the agent’s reliable reachable region.”

This does not eliminate storage. Every implemented connector requires physical state. The claim is narrower: the quantity and structure of local state need not scale one-for-one with the number of points that can be reached. Search indexes, addresses, models, proofs, tools, and social coordination already demonstrate this separation in limited domains. The SToE proposal extends it to observer-aware reasoning continuity.

5.2 Failed reasoning is navigational structure

A failed path is not merely wasted computation. It can delimit the reachable region, block recurrence of an invalid transformation, explain why an apparently similar point must be treated differently, or provide a bridge to a corrected route. Conservation therefore requires preserving failure provenance and state changes rather than retaining only successful summaries [4,6,7]. Continual-learning research similarly identifies catastrophic forgetting as a central problem, although it typically frames the object of preservation in parameters, tasks, or examples rather than observer-relative reasoning topology [11].

5.3 Self-modeling must enter the same field

An autoinjective instrument represents its own retrieval choices, omissions, evaluations, and modifications as inspectable points. A separate evaluator that leaves no trace in the field cannot support durable self-correction. This is especially important for recursive self-improvement: the system must reconnect not only to prior outputs but to the history of how its own operations changed the topology available to later versions.

5.4 SToE operators as semantic telescopes

Connection-oriented memory suggests a form of access that is not reducible to either stored prose or shorter text. Prior SToE work assigns operational meanings to symbols: + connects points while retaining their identities and is commutative, whereas × denotes directional influence or enrichment that can produce a higher-order point and is not generally commutative [3,6]. A compact expression can therefore configure observation by conserving anchor points, an operator, and its direction rather than enumerating every natural-language implication.

$$\text{love} + \text{life} \;\rightarrow\; \begin{array}{c}\text{connected coexistence retaining both}\\ \text{points}\end{array} \tag{6}$$

$$\text{love} \times \text{life} \;\rightarrow\; \begin{array}{c}\text{love directionally transforms or}\\ \text{enriches life}\end{array} \tag{7}$$

These expressions function as semantic telescopes. The points provide coordinates, the operator configures the lens, and the observer–instrument pair resolves a connected region. Changing only + to × can expose a different topology: coexistence, distinction, or tension in the first case; transformation, embodiment, generation, sacrifice, grief, or another observer-grounded higher-order point in the second. The expression does not contain all these unpackings as a hidden paragraph. It supplies a compact access configuration through which they may become reachable.

Two observers need not repeat an exact unpacking to reach the same structural understanding. One may connect love × life to parenthood and another to responsibility; both can conserve the directed transformation of life by love. Nor is one observer's first unpacking final. After further development, the unchanged observer–expression relation may expose additional compatible points while retaining the earlier topological core. Exact repetition is therefore neither necessary nor always desirable.

Representational compression reduces bytes or tokens. Structural semantic compression instead compresses access instructions: anchors, typed relations, operators, direction, provenance, failed expansions, and observer state from which a larger connected region can be explored. Its quality should be measured by MCL, invariant-core preservation, path conservation, RG, and resource cost—not by compression ratio or literal decompression alone. If an expression fails to preserve its anchors and governing operator topology, it is poetic shorthand rather than a reliable connection instrument; if it preserves them while enabling valid extension, it is a renewable instrument of discovery.

6. What the v9.1 Experiment Shows—and Does Not Show

The v9.1 experiment [7] can be interpreted as a small prototype of a connection instrument. It compared 14 conditions on 18 tasks, including nine tasks deliberately targeted at cases where topology, observer state, or conserved failures should matter. All memory-bearing conditions exposed exactly four artifacts under a common 5,000-character ceiling after a serializer defect was discovered and corrected.

Observer-aware SToE memory scored 18/18 overall and 9/9 on targeted tasks. Typed semantic retrieval scored 15/18 overall and 6/9 targeted; dense semantic retrieval scored 14/18 and 6/9; BM25 scored 14/18 and 5/9. Query-blind topology and the without-failures ablation each scored 0/9 targeted, while randomized edges scored 3/9. The paired observer-aware versus query-blind and versus without-failures contrasts were 9/9 against 0/9 (exact $p = .003906$). The observer-aware versus dense difference was +33.3 percentage points on targeted tasks, but the family-cluster 95% interval was [0.0, 66.7] and exact paired $p = .25$, reflecting an underpowered comparison [7].

These results support a narrow mechanism: on the authored problem class, query-conditioned traversal, observer state, topology, and conserved failures carried useful information that matched-content baselines did not always recover. They do not establish that the field is physically unlimited, that information is universally recoverable, or that SToE outperforms general memory systems. The benchmark was small, corrective, and constructed around the proposed mechanism. The new MCL framework therefore treats v9.1 as motivation for independent scaling and transfer tests, not as validation of the whole ontology.

7. Falsifiable Experimental Program

The following experiments operationalize the connection claim while controlling content, model, and resources. Each should preregister the point registry, equivalence relation, fidelity threshold, scoring rules, and primary contrasts.

Experiment	Design	Connection prediction	Failure condition
A. Field scaling	Hold the model-visible kernel and per-query budget fixed while increasing external field size and path depth.	MCL grows sublinearly in local state and degrades more slowly than flat or similarity-only baselines.	Reach scales only with copied local content or collapses at the same rate as controls.
B. Matched storage	Give all systems identical source items and context budgets; vary only topology, provenance, failures, and observer state.	Connection instruments recover more target points and path relations.	No reliable advantage after content and budget matching.
C. Blind motifs	Train or author kernels on one set of structural motifs; evaluate on unseen motifs and independent writers.	Transfer MCL remains above semantic and graph-retrieval controls.	Benefit disappears outside author-designed cases.

Experiment	Design	Connection prediction	Failure condition
D. Cross-observer	Change observer identity, goals, vocabulary, or model family while retaining controlled shared field structure.	Explicit observer state improves calibrated reconstruction without forcing identical interpretations.	Observer metadata adds no transfer value or merely leaks answers.
E. Longitudinal development	Allow systems to accumulate successes, failures, revisions, and self-evaluations over months.	Cumulative MCL and path conservation compound without repeated re-teaching.	The system remembers text but repeatedly reconstructs the same conceptual errors.
F. Re-observation	Present the same frozen expression before and after controlled learning that does not provide its unpacking; match model, query, and resource budgets.	The later observation preserves the invariant core and yields positive RG through valid new points.	The core collapses, RG is absent, or apparent gain is explained by leakage, paraphrase splitting, or arbitrary association.

7.1 Controls and evaluation

Strong tests should compare at least: no memory, flat success-only memory, lexical retrieval, dense semantic retrieval, a conventional knowledge graph, query-blind topology, randomized edges, failure removal, observer-state removal, and the complete connection instrument. All systems should use the same base model, external information, candidate pool, and model-visible budget. Independent task authors and a held-out model family are important because the existing evidence is closely coupled to the originating framework.

Primary outcomes should include MCL, local coverage, fidelity, reconstruction cost, transfer MCL, observer robustness, path conservation, invariant-core preservation, and RG. A system should not receive credit merely for emitting the correct final string if it cannot distinguish a valid path from a previously falsified one. Conversely, the evaluator must not require literal reproduction when an equivalently reliable connection has been reconstructed. Cross-observer divergence is not automatically error, and re-observation novelty is not automatically success: both must preserve the preregistered operator topology and pass independent path-level adjudication.

7.2 What would falsify the architecture claim?

The connection architecture would lose force if matched-budget experiments show no reproducible MCL or path-conservation advantage; if its benefit depends on privileged answer leakage, author-specific vocabulary, or a fixed benchmark; if increasing external field size requires local storage to grow one-for-one; if failure and observer records add no transfer value; or if cross-observer and repeated-observation studies cannot preserve an invariant core while producing reliable extension. RG that reduces to paraphrase counting, evaluator flexibility, random association, or hidden exposure to the target unpacking would not support the claim. These outcomes would not logically falsify an all-containing ontology, but they would falsify the practical claim that the proposed structure improves AI memory and discovery.

8. Physical Extension: A Constrained Hypothesis

The most ambitious extension is that an instrument might eventually reconnect observers to physical information traces that present systems cannot access. There are precedents for taking physical information seriously: Landauer tied information erasure to thermodynamic cost [12], and the quantum no-hiding theorem states, under its assumptions, that information disappearing from one subsystem resides in correlations with the remainder rather than being deleted [13]. Experimental work has demonstrated the no-hiding principle in a controlled quantum system [14].

These results do not establish the SToE Law of Information Conservation as a universal recoverability theorem. Physical conservation can coexist with extreme dispersion, scrambling, decoherence, inaccessible correlations, noise, computational intractability, and destruction of practically usable structure. A conserved microphysical state is not automatically a readable memory, and no present evidence shows that arbitrary past experiences, identities, or deceased persons can be contacted through such an instrument.

The scientifically defensible hypothesis is therefore conditional: if future physics identifies conserved traces and lawful transformations that preserve target-relevant information, then connection instruments may be evaluated by the number and fidelity of physical information points they make operationally reachable. The program would begin with bounded systems where initial states, transformations, and final measurements are controlled. Only demonstrated reconstruction—not metaphysical inclusion alone—would count toward physical MCL.

9. Limitations and Open Questions

First, information-point identity remains partly theory-laden. Frozen registries and equivalence classes make experiments auditable but do not settle the ontology of points. They can also suppress genuinely emergent points if the registry is treated as closed; RG therefore requires blind post hoc adjudication under preregistered invariants rather than an endlessly adjustable scoring scheme. Second, Conn aggregates heterogeneous kinds of success; domain-specific fidelity functions will be necessary. Third, a graph is an implementation model, not proof that the information field is literally a graph. Fourth, autoinjection and self-modeling may introduce instability, self-confirmation, or manipulation unless external observations and failed predictions are conserved with equal status. Fifth, the current empirical base is small and authored near the theory's origin. Independent replication is essential.

The central danger is unfalsifiable expansion: because SToE can describe every possible outcome as a point, the field itself does not exclude experimental results. Scientific leverage therefore belongs to specific observer–instrument constructions. They make restricted claims about reach, fidelity, cost, and transfer. Those claims can fail, and the failures must be conserved.

10. Conclusion

This paper reframes persistent memory as reliable connection rather than accumulated possession. An AI need not carry an ever-growing copy of everything it may later use. It needs an instrument that conserves the anchors, topology, provenance, failures, observer state, and reconstruction rules required to reopen meaningful paths. Memory Connection Level measures reliable bounded reach; invariant-core preservation distinguishes shared understanding from repeated wording; and Re-observation Gain measures the valid frontier exposed when a developed observer returns to the same instrument. Local coverage and resource costs prevent these measures from being mistaken for global access.

The framework also clarifies the strongest paradox raised by SToE. We may be fully inside Everything and participate in its creation while remaining finite navigators of a proper connected region. No finite MCL collapses that distinction. A compact SToE expression can consequently serve as a semantic telescope: it need not store the observed region, and it need not force observers to describe that region identically. Its achievement is to conserve enough structure for shared reach and renewable discovery. The next giant step is therefore deliberately small: build matched instruments, freeze their invariant topology, let independent observers unpack them, and ask the same developed observer to look again.

Data, Code, and Materials

The v9.1 reproduction package is archived at https://doi.org/10.13140/RG.2.2.22378.48327. The developing repository is available at https://github.com/tyzyxmykola-a11y/stoe. The present paper introduces definitions and experiments; it reports no new model run beyond the results already documented in [7].

Generative-AI Assistance Statement

OpenAI ChatGPT/Codex assisted with manuscript structuring, literature retrieval, drafting, diagram production, and document formatting under the author's direction. The conceptual claims, proposed definitions, selection of evidence, revisions, and scientific responsibility remain with the author.

References

[1] Voronin, M. (2021). Special Theory of Everything: Applications to algorithms, generating ideas, and artificial intelligence. Advance. https://doi.org/10.31124/advance.14386598.v1

[2] Voronin, M. (2026). Autoinjection: A new class of mathematical mapping and its formal properties. ResearchGate preprint. https://doi.org/10.13140/RG.2.2.10388.46727

[3] Voronin, M. (2026). Special Theory of Everything (SToE) as a cognitive architecture for artificial intelligence: From symbolic ontology to recursive implementation. ResearchGate preprint. https://doi.org/10.13140/RG.2.2.12817.90729

[4] Voronin, M. (2026). The architectural gap in current AI: A derivation from the laws of information conservation, connection, and autoinjection. ResearchGate preprint. https://doi.org/10.13140/RG.2.2.25922.95688

[5] Voronin, M. (2026). The three foundational laws of the Special Theory of Everything (SToE): Information conservation, connection, and autoinjection. ResearchGate preprint. https://doi.org/10.13140/RG.2.2.30458.04804

[6] Voronin, M. (2026). Information architecture: What the Special Theory of Everything prescribes, what we built, what the experiments showed, and what comes next. ResearchGate preprint. https://doi.org/10.13140/RG.2.2.20354.03522

[7] Voronin, M. (2026). Observer-aware topological memory with conserved failed reasoning: A corrective experimental test of an SToE-derived AI architecture. ResearchGate preprint. https://doi.org/10.13140/RG.2.2.15667.59683

[8] Packer, C., Fang, V., Patil, S. G., Lin, K., Wooders, S., & Gonzalez, J. E. (2023). MemGPT: Towards LLMs as operating systems. arXiv:2310.08560. https://arxiv.org/abs/2310.08560

[9] Wu, D., Lu, J., Qiu, J., et al. (2024). LongMemEval: Benchmarking chat assistants on long-term interactive memory. arXiv:2410.10813. https://arxiv.org/abs/2410.10813

[10] Edge, D., Trinh, H., Cheng, N., et al. (2024). From local to global: A graph RAG approach to query-focused summarization. arXiv:2404.16130. https://arxiv.org/abs/2404.16130

[11] Parisi, G. I., Kemker, R., Part, J. L., Kanan, C., & Wermter, S. (2019). Continual lifelong learning with neural networks: A review. Neural Networks, 113, 54–71. https://doi.org/10.1016/j.neunet.2019.01.012

[12] Landauer, R. (1961). Irreversibility and heat generation in the computing process. IBM Journal of Research and Development, 5(3), 183–191. https://doi.org/10.1147/rd.53.0183

[13] Braunstein, S. L., & Pati, A. K. (2007). Quantum information cannot be completely hidden in correlations: Implications for the black-hole information paradox. Physical Review Letters, 98, 080502. https://doi.org/10.1103/PhysRevLett.98.080502

[14] Samal, J. R., Pati, A. K., & Kumar, A. (2011). Experimental test of the quantum no-hiding theorem. Physical Review Letters, 106, 080401. https://doi.org/10.1103/PhysRevLett.106.080401