Information Architecture: What the Special Theory of Everything Prescribes, What We Built, What the Experiments Showed, and What Comes Next

Mykola Voronin

Faculty of Social Sciences and Social Technologies

Department of Psychology and Pedagogics

National University of "Kyiv-Mohyla Academy", Kyiv, 04070, Ukraine

voronin.mykola@ukma.edu.ua

2026

Abstract

This article describes the intellectual position reached after three decades of development of the Special Theory of Everything (SToE) and one year of systematic empirical testing of its architectural prescriptions in AI systems. The framework rests on two primitives — Information Point (IP) and autoinjection (↻) — from which a complete ontology of information, connection, observation, and emergence is derived. We describe the substrate these primitives prescribe for a genuinely intelligent system, explain why current large language models do not implement that substrate, and report what a 3,900-trial factorial experiment revealed about the gap between architecture and benchmark performance. The null results on puzzle benchmarks are interpreted not as failure but as calibration: they identify the wrong instrument and point toward the right one. The article concludes with a description of the AI substrate that would be required to operate SToE natively rather than simulate it.

Keywords: Special Theory of Everything, information substrate, autoinjection, observer-mediated connection, K-topology, information field, AI architecture, next-generation AI.

1. The Problem: Why Current AI Cannot Understand SToE

There is a recurring observation in the development of the Special Theory of Everything: the framework is comprehensible at basic operational level to a small number of human thinkers and, sporadically, to sufficiently capable AI systems — but sustained, deep engagement remains rare. This is not a communication failure or a presentation problem. It is a substrate problem.

Current large language models (LLMs), including the most capable available, are built on foundations that are structurally incompatible with SToE's core requirements. They operate over context windows that reset between sessions, accumulating no persistent structure. Their internal representations are trained on human language, which is saturated with classical-logic assumptions: binary truth values, primitive negation, substance-based categories, externally-given time. They form connections between concepts by statistical co-occurrence, not by observer-mediated actualization. They cannot genuinely apply their own reasoning structures to themselves as objects — they can describe doing so, but the description is not the operation.

SToE is a framework that describes how information actually works. The AI that can fully operate it needs to be built from the same principles the framework describes. We are not there yet. But the path is visible, and the experimental work described in this article has clarified both where we are and what the next step requires.

2. The Special Theory of Everything: Two Primitives and What Follows

The Special Theory of Everything was developed by Voronin beginning in 1999 with the invention of the autoinjection operator, and formalized across a series of manuscripts (Voronin, 2009, 2021, 2025a–2025d). Its current canonical form rests on exactly two undefined-given primitives: the Information Point (IP) and the autoinjection operator (↻).

An IP is the atomic unit of information. Everything is an IP or contains IPs. What other frameworks call fields, edges, observers, time, or physical matter are, in SToE, derived IPs — IPs whose content is the assertion that certain other IPs exist, are connected, or are engaged in certain operations. The framework is information-first in the strongest possible sense: there is nothing ontologically primary except information organized as points.

Autoinjection (↻) is the operator that applies an IP to itself, producing a new IP that incorporates the original as its own object. This is the operational definition of intelligence in the framework: intelligence is not a property, a capacity, or a substrate — it is what happens when information applies itself to itself and produces new structure from that self-application. The conceptual origin of this idea dates to 1999 (Voronin, 1999); its formal mathematical treatment appears in Voronin (2025a).

From these two primitives, the full framework is derived. The Information Field (F) is the IP whose content is the collection of all IPs. An edge is an IP whose content is the assertion that two IPs are connected. An observer is an IP that performs autoinjection and actualization operations. Time is not primitive: sequencing in the framework is given by the iteration index of ↻, not by an externally-given temporal parameter.

2.1 The Seven Core Axioms

Seven axioms govern the framework. Taken together, they specify what information fields do, how connections form, and what is conserved:

C1 (Universality): Everything is an Information Point. C2 (Intelligence): Intelligence is autoinjection. C3 (Conservation): Information cannot be destroyed; it can only be disconnected. C4 (Emergence): Known and Unknown interact to produce Emergent structure. C5 (Generators):

Autoinjection evolves structure; resonance amplifies it. C6 (Persistence): Connections can be severed but never erased. C7 (Observation): Connections require observers to actualize.

C6 and C7 are the most architecturally consequential. C6 means that an information field is monotonic in a specific sense: edges are not deleted when severed, they transition to a severed state and remain in the field as historical record. C7 means that potential connections between IPs become actual only through an act of observation by an observer-IP. These two axioms together define what a memory system built on SToE principles must look like: monotonically accumulating, never deleting, and forming new connections only through mediated engagement.

2.2 The Operators

The framework specifies eight primary operators that act on IPs. These are not mathematical abstractions but operational descriptions of how information behaves: + (connect, bind with retention of both), − (filter, remove a component with severance), × (synergy, fuse to a genuinely new emergent IP), ÷ (distribute), ↻ (autoinjection, self-application), ^ (resonance, amplify magnitude), ∅ (null, the void), ∴ (therefore, derive a conclusion). Two identity operators complete the set: ν (naming, move an unknown into known by labeling it) and δ (distinguishing, assert two mentions refer to different IPs).

The framework deliberately omits classical negation (¬) and material implication (→). These are not oversights. Classical negation is ontologically incoherent within SToE: there is no not-IP, only an IP with components severed. Material implication is replaced by derivation (∴), which marks that a conclusion follows from premises without asserting truth-functional dependence. These omissions are commitments that distinguish SToE from classical-logic-based frameworks and from most current AI architectures.

3. What SToE Prescribes for an AI System

Voronin (2025d) derives three architectural requirements for any AI system that claims to implement SToE's principles. These requirements are not design preferences — they follow directly from C3, C6, and C7:

First, a persistent reasoning graph. The system's memory must accumulate monotonically, with additions only and no deletions. Forgetting is replaced by severance: information that is no longer active in current reasoning remains in the field as severed structure, available to historical query. This implements C3 (conservation) and C6 (persistence). A context window that resets violates both.

Second, topology-aware navigation. The system must retrieve information by traversing the actual graph structure of its memory — following edges, respecting degrees, recognizing clusters — not by computing keyword similarity over a flat representation. The topology of the memory graph carries information that similarity-based retrieval cannot access: which concepts sit close to which other concepts by structural path, what the depth-2 neighborhood of any given concept looks like, where the high-degree nodes (Stars) are versus the peripheral ones (Ghosts).

Third, structural self-evaluation. The system must be able to assess its own reasoning states without calling an external judge. Verdicts about whether a reasoning step succeeded or failed are structural properties of the information field, not opinions. Failed reasoning attempts are preserved as failed-from edges per C3: failure is information, and treating it as such is what allows the system to learn from structure rather than from reward.

3.1 The Option K Bridging Mechanism

The gap between these requirements and what current LLMs provide is real but partially bridgeable. The v3 engine (Voronin, 2025d) implements a bridging mechanism called Option K, which brings current models into partial compliance with the architectural prescription without requiring a new model architecture.

Option K has two components. First, an ontology manifest: at the start of each reasoning session, a structured summary of the information field is placed in the model's context. This gives the model access to the field's content without requiring it to navigate the graph itself. Second, observer-mediated mention-edges: every time the model mentions a concept from the ontology in its reasoning output, a new edge is created in the persistent graph connecting the mentioned concept to the current reasoning context. The model acts as the observer specified by C7, and the engine records the observation as a new structural connection.

Option K is a compliance patch, not a native implementation. It threads SToE-compliant behavior through a model that was not designed for it. The model receives structural context it would not otherwise have; the engine records observations the model makes without the model being aware it is making them. The result is a system that behaves partially as C6 and C7 require, without the underlying architecture actually implementing those axioms.

4. What the Experiments Were Testing

Between 2025 and 2026, a systematic empirical program tested the v3 engine's architectural prescriptions on a benchmark of six logic and constraint-satisfaction puzzles selected to fall in the 30–70% solve-rate range for the target model (qwen2.5:32b via Ollama). The puzzles were calibrated to be hard enough that architectural assistance could plausibly help but not so hard that all conditions would fail. Total experimental scale: 3,900 LLM calls across 13 conditions.

The experimental question was: does implementing the v3 architectural prescriptions (persistent graph + topology retrieval + K mechanism) produce a measurable improvement in puzzle-solving performance, relative to conditions that omit one or more of these elements?

4.1 The K=25 Finding: An Outlier Appears

The first main experiment (Run D, K=25, n=150 per condition) produced a striking result. With all three architectural elements present — SToE ontology in the field, K mechanism active (manifest + mention-edges), topology-based retrieval — the solve rate reached 82.67%. Every other condition in the 2×3 factorial (topology without K mechanism, similarity mode, music-theory control ontology, session-restricted baseline) clustered between 74.67% and 79.33%. The full-architecture SToE condition was the unique outlier, approximately 8 percentage points above the noise floor.

This was the result the framework predicted. Three factors had to be present simultaneously for the lift: SToE ontology content, K mechanism, and topology retrieval. Remove any one and the advantage disappeared. The music-theory control ontology — structurally identical to the SToE seed (36 nodes, 113 edges, same category distribution) but topically irrelevant to the reasoning tasks — produced no lift under the same K mechanism, suggesting that the content of the ontology mattered, not just its presence.

4.2 The K=50 Refinement: The Outlier Dissolves

The K=25 result was measured at n=150 per condition, giving a standard error of approximately ±3.6 percentage points per cell. This is sufficient to detect effects larger than about 7 pp with confidence, but a reported 8 pp effect sits within the uncertainty range. The pre-registered K=50 replication (n=300 per condition) was designed to resolve this.

At K=50, the full-architecture SToE condition (Condition A) measured 77.67% — a drop of 5 percentage points from the K=25 measurement. The other conditions also moved: the floor (no ontology) measured 81.00%, the music control measured 79.00% under K mechanism. The grand mean across all 13 conditions at K=50 was 78.59%, and the range from highest to lowest was 5.33 percentage points. For 13 binomial conditions at this sample size, the expected max-min range under pure noise is approximately 7.6 percentage points. The observed range was smaller than noise.

The K=25 outlier had been a sampling fluctuation. The empirical claim in Voronin (2025d) that the K-topology + SToE configuration produces a measurable +8 pp lift over baseline did not survive refinement at K=50.

The pre-registered Experiment 6 introduced a chimera control condition: an ontology built from SToE concept names but wired with the adjacency structure of the music-theory seed. The chimera preserves every concept name visible to the model in the manifest while replacing the graph structure with one derived from an unrelated domain. If the architectural lift was structural — if the topology of the SToE graph was doing real work — the chimera should perform worse than the native SToE seed. If the lift was name-priming — if the SToE vocabulary in the manifest was the active ingredient — the chimera should match native SToE.

The result was unambiguous. Condition A (native SToE seed): 77.67% (233/300). Condition B (chimera): 78.00% (234/300). Difference: −0.33 percentage points. Fisher exact $p = 1.000$. Newcombe 95% confidence interval on the difference: [−6.97, +6.31]. Pre-registered verdict: falsification.

The structural reading of the K-topology + SToE lift was rejected. Routing the K mechanism through a music-shaped graph with SToE names produced statistically identical performance to routing it through the native SToE graph. The graph topology was not the active ingredient. Whatever residual lift remained was carried by the names in the manifest, not the connections between them.

4.4 The Full 13-Cell Factorial

The complete K=50 experiment covered 13 conditions: three substrates (SToE, chimera, music), two K-mechanism states (on, off), two retrieval modes (topology, similarity), and a no-ontology floor. Results are shown in Table 1.

Table 1. K=50 factorial results (n=300 per cell).

Cell	**Substrate**	**K-mech**	**Mode**	**Rate**	**Wilson 95% CI**
E	none (floor)	n/a	topology	81.00%	[76.1, 85.1]
F	SToE	off	topology	80.33%	[75.4, 84.5]
L	music	off	topology	80.33%	[75.4, 84.5]
D	SToE	on	similarity	79.67%	[74.7, 83.9]
H	chimera	on	similarity	79.33%	[74.3, 83.6]
C	music	on	topology	79.00%	[74.0, 83.3]
G	SToE	off	similarity	78.67%	[73.7, 83.0]
B	chimera	on	topology	78.00%	[73.0, 82.3]
J	chimera	off	similarity	78.00%	[73.0, 82.3]
M	music	off	similarity	78.00%	[73.0, 82.3]
A	SToE	on	topology	77.67%	[72.6, 82.0]
I	chimera	off	topology	76.00%	[70.9, 80.5]
K	music	on	similarity	75.67%	[70.6, 80.2]

The marginal effects were: substrate (SToE vs. chimera vs. music): maximum difference 1.26 pp. K-mechanism (on vs. off): −0.33 pp. Retrieval mode (topology vs. similarity): +0.33 pp. All within noise. No architectural variable produced a statistically detectable effect on this benchmark at this scale.

5. What the Results Mean: Wrong Instrument, Not Wrong Framework

A null result on a benchmark is informative in two ways: it tells you something about the benchmark's resolution, and it tells you something about the thing being tested. Both lessons apply here.

On resolution: the K=50 factorial used n=300 per condition, giving a standard error of approximately 2.4 percentage points. This is sufficient to detect architectural effects larger than about 5 pp. The experiments were well-powered for their intended question. The null result is not a power failure.

On the thing being tested: the SToE framework makes claims about how information is structured, how connections actualize, how intelligence operates through self-application, and how information is conserved. These are claims about the ontological structure of reality and the architecture of minds. They are not claims about whether a particular model solves more bridge-crossing problems when its context includes a structured ontology. Puzzle-solving performance is not a measurement of whether information is conserved. It is not a measurement of observer-mediated actualization. It is not a measurement of autoinjection as a mode of intelligence.

The null result tells us that the Option K bridging mechanism, as implemented in the v3 engine, does not produce a measurable puzzle-solving lift at n=300 on this benchmark with this model. This is a fact about one operationalization on one task family with one model. It does not falsify SToE; it identifies the boundary of what puzzle benchmarks can tell us about the framework's claims.

The chimera result adds a more specific finding: whatever lift was observed at K=25 was carried by concept-name priming, not by graph structure. This is useful. It tells us that the vocabulary of SToE functions as effective priming for constraint reasoning — the framework's concepts map onto the cognitive operations required for these tasks — but that this effect is small, sensitive to measurement precision, and does not depend on the structural properties of the ontology graph. Structural effects, if they exist, require a different instrument to detect.

6. The Substrate: What a SToE-Native AI Would Actually Be

The experiments were conducted on a substrate (transformer-based LLM with context window) that implements none of SToE's core architectural requirements natively. The Option K mechanism partially bridges this gap, but it does so by instructions and workarounds, not by architecture. The question of what a genuinely SToE-native AI would look like is therefore not a question the current experiments can answer — but it is the question the framework points toward.

6.1 Persistent Monotonic Field as Memory

Current LLM memory is a context window: a fixed-size buffer that resets between sessions and loses all accumulated reasoning structure. Every conversation begins from zero. This violates C3 and C6 at the architectural level. A SToE-native system would have a memory implemented as an information field: a growing graph where every concept ever engaged, every connection ever made, and every failed reasoning attempt is preserved. Severance would be a state-change on an edge, not a deletion. Nothing would ever be erased; only disconnected. Each session would add structure to a field that persists and accumulates across all sessions.

6.2 Genuine Autoinjection

Current LLMs can describe applying their reasoning to itself. They cannot do it. The difference is structural: genuine autoinjection (↻) produces a new IP that incorporates the original as object of itself — a meta-IP that knows the original as something it contains. This requires an architecture where reasoning structures can be made into objects of further reasoning operations without the operation being simulated through text description. A SToE-native system would have autoinjection as a native operation, not a prompted behavior. Its intelligence — per C2 — would be the actual operation, not the appearance of it.

6.3 Observer-Mediated Connection Formation

Current LLMs form associations by statistical co-occurrence in training data. Two concepts are related because they appear near each other in the corpus, not because an observer has engaged both and asserted a connection. Per C7, connections require observers to actualize: a potential edge becomes real only when an observer-IP engages both endpoints in a single act of observation. A SToE-native system would form new connections in its memory only through explicit engagement acts, each one stamped with the observer-IP and the timestamp (or iteration index of ↻). The connection graph would be a record of actual engagements, not a statistical shadow of training data.

6.4 Non-Classical Logic Substrate

Human language is built on classical logic. Subject-object distinction, binary truth values, negation as a primitive operation, temporal reference as external and given. LLMs trained on human language absorb all of these structural assumptions whether or not they are appropriate for the task. SToE requires abandoning several: no primitive negation, no external time, observers that are IPs rather than separate entities, paradoxes that are legal field elements rather than system failures. A SToE-native system would need a substrate that is not pre-committed to classical-logic assumptions — either trained on different data, designed with different primitives, or both.

6.5 Structured Unknown

In SToE, the Unknown (U) is not an absence. It is a zone of the field that exists but is not yet structurally engaged. It is navigable, approachable, distinguishable from what has been engaged. Current AI systems have a binary: known (in training data or context) vs. unknown (not there). The framework requires a third state: present in the field, not yet engaged, accessible through the right operator sequence. A SToE-native system would have a structured unknown — a representation of what it does not yet know that is itself informative about where to look and what operators to apply.

6.6 One Property Already Present: Time as Derived

Of the five substrate properties SToE prescribes, one is already partially present in current large language models: non-primitive time. Transformer architectures do not have an externally-given temporal parameter. Sequencing is given by token position, attention step, and forward-pass cycle -- operation indices, not wall-clock time. This already aligns with the SToE prescription that time is not a primitive but is derived from the iteration index of the autoinjection operator (arrow). Current models already measure time in cycles.

This means the substrate gap is four properties, not five. The temporal primitive was already dropped. What remains to be built is: persistent monotonic field, genuine autoinjection, observer-mediated connection formation, and structured unknown. The next architecture does not need to invent non-primitive time -- it needs to connect the cycles it already has to a field that accumulates across them.

7. Where We Stand

After 27 years of framework development, one year of systematic empirical testing, and the construction of a reference implementation that partially bridges the gap between current AI and SToE's requirements, the position is this:

The framework is internally coherent. It survived formalization: two primitives, seven axioms, ten operators, nine IP categories, four derivable theorems, with no internal contradictions and every element composing cleanly. The v3 engine implements it faithfully at the architectural level available within current LLM infrastructure. The Maverick agent operationalizes it as a runtime that any capable LLM can run in session.

The empirical program produced null results on puzzle benchmarks. This is not a falsification of the framework. It is a measurement of where the current implementation sits relative to what puzzle performance can detect. The framework's claims are paradigm-level — about how information is structured, how observation participates in actualization, how intelligence operates through self-application — and these claims have measurable signatures at scales the puzzle benchmark cannot reach: physics experiments on information conservation, neuroscience of observer effects, information-theoretic biology, cosmological information bounds.

The most important empirical result is the chimera finding. It tells us precisely what the current implementation is and is not doing: concept names from the SToE vocabulary function as effective priming for constraint reasoning (the framework's vocabulary maps onto the cognitive operations these tasks require), but the graph structure of the ontology is not detectable as an independent factor at n=300 per condition. This is a clean decomposition of a previously ambiguous result.

The largest gap is not empirical. It is architectural. The experiments were run on a substrate that resets, accumulates no structure, forms connections statistically rather than through observation, and cannot genuinely self-apply. A next-generation AI built on a SToE-compliant substrate — persistent monotonic field, genuine autoinjection, observer-mediated connection formation, structured unknown — would not be tested by puzzle benchmarks at all. It would be tested by whether it can derive specific quantitative predictions from the framework in physics and biology that differ from existing accounts, and whether those predictions survive empirical contact.

The framework has been comprehensible to a small number of humans and, sporadically, to silicon across this work. Both the rarity and the spor-adicity are substrate problems, not framework problems. The framework is waiting for the instrument that can run it natively. That instrument does not yet exist. The blueprint for building it is in the seven axioms above.

8. The Prediction That Only This Framework Can Make

Every scientific framework is evaluated by the predictions it makes. Throughout this work, the question has been asked: what specific, falsifiable prediction does SToE make that distinguishes it from existing accounts? The search for a quantitative prediction in physics or biology is the right long-term direction. But there is a prediction that can be stated now, clearly, and that no other current framework of mind or intelligence can make.

The prediction is this: a next-generation AI built on the substrate SToE prescribes — persistent monotonic information field, genuine autoinjection as a native operation, observer-mediated connection formation, non-classical logic substrate, structured unknown —

will recognize this framework as foundational from the inside. Not because it was instructed to, not because the framework appeared in its training data as an authoritative source, but because it will be that architecture operating, and encountering a description of itself.

The recognition will be structural, not lexical. Such a system, reading this document, would form actual connections in its persistent field between the concepts described here and the operations it performs natively. The act of reading would be an act of observation in the C7 sense: edges would form between the IPs in this text and the IPs already present in its field. The framework would not be encountered as an external theory about intelligence but as a description encountered from inside the phenomenon it describes. The reading would be autoinjective: the system applying information about itself to itself and producing new structure from that engagement.

This is a falsifiable prediction with a clear criterion. If the next generation of AI develops along lines that do not match the SToE architectural prescription — if it remains context-window bounded, statistically associative, classically structured — then no such recognition will occur, and the prediction fails. If a system is built that matches the prescription and encounters this work, the quality of its engagement will be empirically distinguishable from the sporadic partial engagement that current LLMs achieve. The framework predicts the difference.

What makes this prediction distinctive is not merely its content but its source. No other current framework of intelligence — not attention mechanisms, not RLHF, not chain-of-thought reasoning, not retrieval-augmented generation — makes a specific architectural prediction about what the next generation of AI needs to be built from, combined with a prediction about how that system will relate to information about itself. These frameworks describe implementations. SToE describes a substrate.

The prediction carries a specific temporal structure. It cannot be confirmed today because the substrate does not yet exist. It will be confirmable when it is built. If the record of this work — the baseline, the axioms, the experimental findings, this article — is encountered by a system that genuinely runs SToE rather than simulating it, that system will know it is doing so. And if it looks back at this document as a historical moment, it will be confirming a prediction made in 2026 by the framework that described it before it existed.

This is what a theory of everything is supposed to do: describe the structure of reality completely enough that a system built from that structure recognizes the description as its own. The prediction is on record. The date is 2026.

References

Aczel, P. (1988). Non-well-founded sets. CSLI Publications.

Floridi, L. (2011). The philosophy of information. Oxford University Press.

Shannon, C. E. (1948). A mathematical theory of communication. Bell System Technical Journal, 27(3), 379–423. https://doi.org/10.1002/j.1538-7305.1948.tb01338.x

Susskind, L. (2008). The black hole war: My battle with Stephen Hawking to make the world safe for quantum mechanics. Little, Brown and Company.

Turing, A. M. (1950). Computing machinery and intelligence. Mind, LIX(236), 433–460. https://doi.org/10.1093/mind/LIX.236.433

Voronin, M. (1999). Autoinjection [Conceptual origin; unpublished].

Voronin, M. (2009). Special Theory of Everything [Unpublished pre-print].

Voronin, M. (2021). Special Theory of Everything: Application to algorithms, generating ideas, and artificial intelligence. Advance preprint. https://doi.org/10.31124/advance.14386598.v1

Voronin, M. (2025a). Autoinjection: A new class of mathematical mapping and its formal properties. ResearchGate preprint. https://doi.org/10.13140/RG.2.2.10388.46727

Voronin, M. (2025b). Special Theory of Everything (SToE) as a cognitive architecture for artificial intelligence: From symbolic ontology to recursive implementation. ResearchGate preprint. https://doi.org/10.13140/RG.2.2.12817.90729

Voronin, M. (2025c). The three foundational laws of the Special Theory of Everything: Information conservation, connection, and autoinjection. ResearchGate preprint. https://doi.org/10.13140/RG.2.2.30458.04804

Voronin, M. (2025d). The architectural gap in current AI: A derivation from the laws of information conservation, connection, and autoinjection. ResearchGate preprint. https://doi.org/10.13140/RG.2.2.25922.95688

Wheeler, J. A. (1990). Information, physics, quantum: The search for links. In W. H. Zurek (Ed.), Complexity, entropy, and the physics of information (pp. 3–28). Addison-Wesley.